

Static Scheduling in Clouds

Thomas A. Henzinger Anmol V. Singh Vasu Singh
Thomas Wies Damien Zufferey

IST Austria

ExCape meeting, June 10 2013, Berkley

Motivation (1)

Cloud computing gives the *illusion* of ∞ (virtual) resources.

Actually there is a finite amount of (physical) resources.

We would like to efficiently share those resources:

- ① being able to distinguish high priority (serving customer *now*) from low priority (batch) requests;
- ② schedule accordingly.

Therefore, we should be able to **plan ahead** computations.

Dynamic Scheduling: use work queues, priorities, but limited.

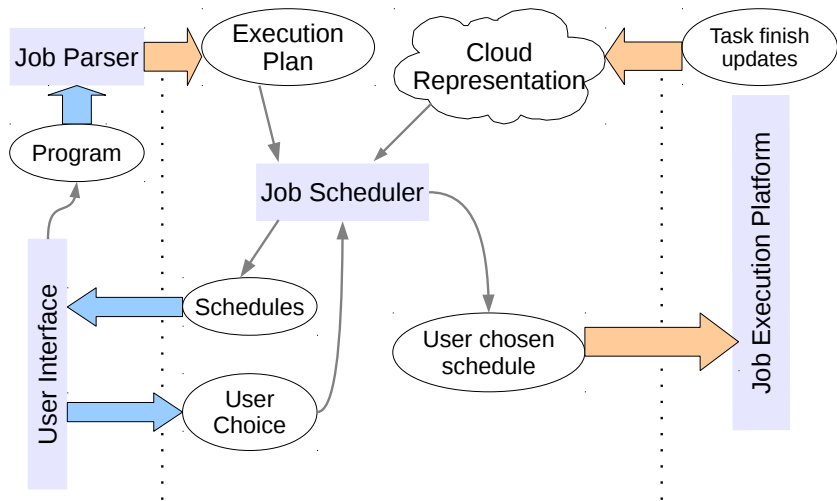
Without knowledge of jobs, this is the best you can do.

We need to ask the user for:

- what kind of resources his job require;
- a deadline/priority for his job.

In exchange we can give him an expected completion time.

We can also offer choice. (time is money.)




Giving incentive to plan in advance

The scheduler returns not one but many possible schedules with different finish times.

Use a pricing model to associate a cost to the schedules.


Include the “scheduling difficulty” in the cost, give a discount to schedule with later finish time.



Problem: static scheduling is *hard*.

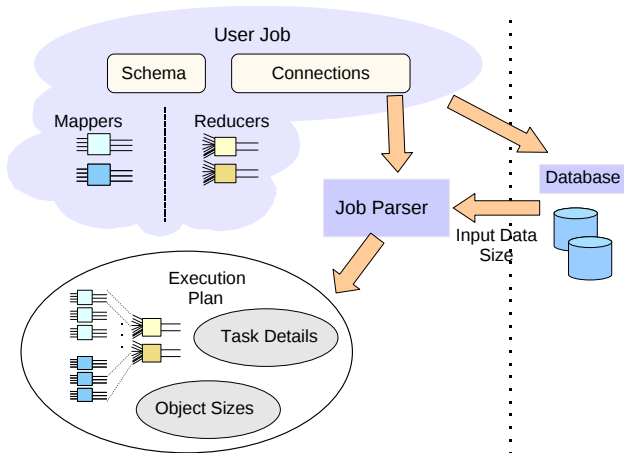
Only possible if the scheduler can handle the work load.


So we set up to make scheduling cheap(er).



- A Job is a directed acyclic task (DAG) of tasks.
- Node are marked with **worst case duration**.
- Edges are marked with **data transfer**.
- duration and data can be parametric in the input.

Parametric Jobs





Datacenter as a tree-like graph:

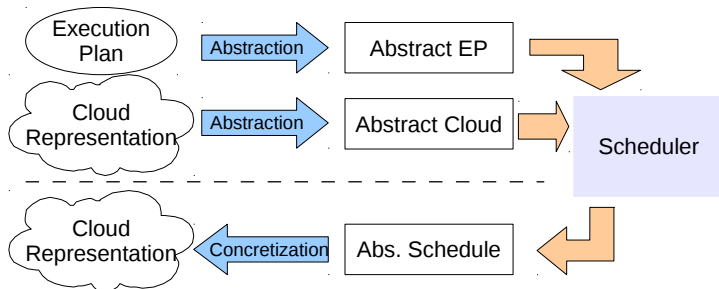
- internal nodes are router;
- leaves are compute nodes (computation speed);
- edges specifies the bandwidth.

Scheduling Large Jobs using Abstraction [EuroSys 2011]

Assumption: job and infrastructure **regularity**


Idea: regularity makes large scale scheduling feasible

How: Using abstraction techniques




Abstraction for jobs:

Group independent tasks as per a topological sort. Merge them into an abstract task.




Abstraction for jobs:

Group independent tasks as per a topological sort. Merge them into an abstract task.




Abstraction for jobs:

Group independent tasks as per a topological sort. Merge them into an abstract task.




Abstraction for jobs:

Group independent tasks as per a topological sort. Merge them into an abstract task.




Abstraction for jobs:

Group independent tasks as per a topological sort. Merge them into an abstract task.




Abstraction for infrastructure:

Merge nodes to according to network topology:




Abstraction for infrastructure:

Merge nodes to according to network topology:




Abstraction for infrastructure:

Merge nodes to according to network topology:



Experiments, part 1: simulation

datacenter:

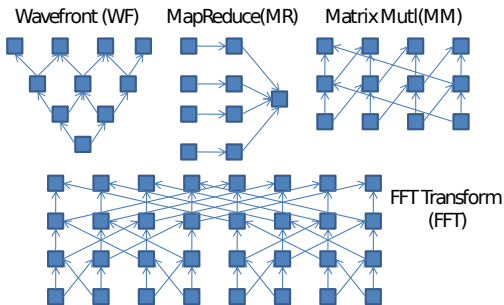


2-tier datacenter

Half of the nodes: speed x ,

Other half: speed $1.5x$


On the job side:




Experiments: the cost of abstraction

We then compare Fisch and Blind to a concrete greedy scheduler (baseline) on a sequence of 100 jobs (10-5000 tasks each). Latency is given per tasks.


Scheduler	Latency (ms)	Utilization
Baseline	293	96 %
Fisch	0.27	92 %
Blind	0.16	91 %



(a) Baseline





(b) Fisch



(c) Blind

Experiments: scaling

- C1: 2000 nodes,
20 per rack
- C2: 1600 nodes,
40 per rack
- C3: 4000 nodes,
20 per rack
- C4: 8000 nodes,
20 per rack
- C6: 1000 nodes,
500 per rack



Caution: static scheduling alone will not work.

- Task duration are conservative estimates;
- Variability of the performance of the compute node.

We use static scheduling with backfilling.



Job:

- The jobs are MapReduce jobs doing image transformation.
- Mapper: 8.1 seconds on average, estimate is 40 seconds
- Reducer: Identity operation

Infrastructure:

- Hadoop streaming version 0.19.0
- Amazon EC2 m1.xlarge instances (15GB RAM, 4 cores)
- Number of mappers = 50 * number of instances

Experiments: compared to Hadoop



Observations:

- The Hadoop framework requires large runtime overhead: results in slowdown of the job execution.
- Static scheduling allows to prefetch data, whereas dynamic scheduling does not

There is an opportunity to apply methods developed to solve computationally hard problem in verification to other area. While preserving a solid theoretical basis.

Questions ?